NotebookLab

Your thinking partner, on your machine.

A desktop notebook that reads your PDFs, notes and scans, then answers with the page it used.

100%
Offline. Your data never leaves your machine.

0
API keys required. Download, install, start working.

3
Platforms. Windows, macOS, and Linux.

Amey Thakur with Archit Konde
github.com/Amey-Thakur/NotebookLab

Free and open source · MIT

START HERE

What this is, in one page

NotebookLab imports your documents, indexes them on your machine, and answers questions about them with a model that runs there too. Every answer names the document, the heading and the page it drew from, so you can check it rather than trust it.

It is an application, not a service. No account, no key, no upload step. Once a model is downloaded it works with the network cable out.

WHO IT IS FOR

Anyone whose documents are not allowed to leave: unpublished research, files under a compliance rule, a term of scanned handouts, internal runbooks.

WHAT IT IS NOT

> Not a chat wrapper. Retrieval and citation are the product; the model is swappable.

> Not a cloud service with a local mode. There is no server component at all.

> Not the easiest choice if your files are free to travel. If they are, use anything.

THE HONEST LIMITS

A local model is slower than a hosted one and weaker at synthesis across many documents, and it knows nothing past its training data. Cloud providers are supported for exactly those cases, and stay off until you add a key.

WHAT AN ANSWER LOOKS LIKE

What sample size did the 2019 trial use?

1,204 participants, after 87 withdrawals.

trial-report.pdf | Methods | page 14

Open the passage

FROM THE MAKERS

We did not set out to build a product.

We wanted somewhere to read our own files without handing them to anybody.

Amey Thakur

Everything here follows from one rule: if a feature needs your documents to leave the machine, it does not ship.

That rule cost us features. It is also the only reason this is worth downloading.

Archit Konde

Most of my work is invisible when it works: a scan that parses, a bad file that fails without taking the app down.

That is what decides whether you open a tool a second time. If it breaks on your files, tell us. That is worth more to us than a star.

THE PLEDGE

No account
Nothing to sign up for, ever

No telemetry
No analytics, nothing phones home

No upload
Your files stay on your disk

Signed
Every commit, every release artifact

SIGNED

NotebookLab

CERTIFICATE OF AUTHENTICITY

The Makers' Pledge

The tools we think with have quietly become products that watch us. The notebook turned into a subscription, the private note into a data point, and the quiet place to work into something that studies what you do inside it. We wanted our thinking partner back, so we built one that answers only to you.

Every part of NotebookLab was made by the two of us, one considered commit at a time, each with a single clear purpose and read closely before it was trusted. We made it the way we would want a tool made for us.

To everyone who makes it part of their work, we give our word.

Yours alone.
Your notes, documents, and questions never leave this machine.
No accounts, no telemetry, no quiet calls home.

Open always.
Every line is here to be read, questioned, and made better
by anyone who cares to look.

Honest work.
We ship only what we run ourselves, and when we get something
wrong we say so plainly and fix it in the open.

Amey Thakur — @Amey-Thakur

Archit Konde — @Archit-Konde

SIGNED

Signed with our own hands, and with the key that signs every commit and release, so this promise can be checked and not merely trusted.

First released to the public in 2026 · MIT License

SHA256:h+W7L5p+iSLNi8FfG+Svgwtr/shv3l0JsCYq/o6VBis

Click to download the full certificate

The Makers' Pledge

The note on the previous page is what we meant. This is what we signed.

It ships inside the app, under About, and it is signed with the same key that signs every commit and every release artifact. That is the point of it: a promise you can verify rather than one you have to take on faith.

WHAT IT COMMITS US TO

Yours alone. Your notes, documents and questions never leave the machine.

Open always. Every line is there to be read and argued with.

Honest work. We ship what we run ourselves, and say so when we get it wrong.

WHERE TO FIND IT

In the app
About, at the bottom of the page

In the repository
docs/brand/makers-pledge.svg

On the site
Download the certificate, full resolution

Everything it does, on your own machine

Reads your documents, answers with the page it used. Free, open source, no account.

READS

PDF
including scanned pages

Word
.docx

Markdown and text
any plain file

Images
read locally by OCR

Drag and drop
or Open with

ANSWERS WITH

The document
which file it came from

The heading
where inside it

The page
so you can check it

The passage
one click away

Search
full text and semantic

RUNS ON

Bundled llama.cpp
any GGUF model

Ollama
if you already use it

OpenAI-compatible
any endpoint

Anthropic, Gemini
your own key

Nothing at all
works offline

ALSO DOES

Connected notes
with a graph

Document outline
jump anywhere

Prompt Studio
build and reuse

Export
Word and RTF

Command palette
Ctrl+K

Windows, macOS and Linux · updates itself, every release signed · MIT licensed · no account, no telemetry, nothing uploaded

Amey Thakur with Archit Konde
github.com/Amey-Thakur/NotebookLab

Your thinking partner, on your machine.

THE IDEA

An answer without a source is a rumour with good grammar

The hard part of local AI is not the model, it is the retrieval. A small model that cites its source is more useful than a large one that does not, because you can check it in seconds.

So provenance is built into the shape of the data rather than added at the end. Every format parses down to the same chunk: PDF, Word, Markdown, plain text, and images through OCR. That chunk carries which file it came from, which heading it sat under, and which page it was on.

A chunk that does not know its own origin is never created. Retrieval therefore never has to reconstruct where something came from, and no shortcut later can quietly lose it.

WHAT A CITATION GIVES YOU

FIELD	WHAT IT IS FOR
Document	Which file the claim came from
Heading	Where inside that file
Page	So you can open it and read the passage
Passage	One click, the exact text used

Deciding this shape before writing a line of retrieval was the choice that saved the most time on the project. Every shortcut that looked tempting later would have cost provenance, and provenance is the product.

Four people who cannot send documents away

The constraint is the reason this exists. Each of these is a real workflow.

THE RESEARCHER

Sixty papers
one literature review

Ask across all of them
not one at a time

Every claim cited
paper, section, page

Quote it directly
the citation is written

Nothing uploaded
preprints stay yours

LEGAL AND COMPLIANCE

Contracts
that cannot leave

Ask the bundle
not each file

Check the clause
it quotes the source

No network needed
air-gapped if you like

No third party
nothing to review

THE STUDENT

Lecture notes
a whole term

Scanned handouts
OCR reads them

Ask a concept
it points at the slide

Search everything
typed and photographed

Free
no subscription

THE ENGINEER

Design docs
and runbooks

Ask mid-incident
no tab hunting

Get the section
not a paraphrase

Stays on the laptop
internal is internal

Works offline
even on a plane

If your documents are allowed to leave, use whatever is easiest. This is for when they are not.

Amey Thakur with Archit Konde
github.com/Amey-Thakur/NotebookLab

Your thinking partner, on your machine.

Where your documents actually go

Every tool that reads your files needs a copy. The question is whose disk it is on.

YOUR MACHINE

- **Your documents**
  PDF, Word, notes, scans, imported once
- **The index**
  chunks that remember file, heading and page
- **The model**
  runs on your own CPU or GPU
- **The answer**
  with the passage it used attached

A cloud model

Optional, and off unless you add your own key. Turn it on and your text does go to that provider, which is exactly why it is a choice, not a default.

With a local model there is no request to make, so there is nothing to intercept.

Amey Thakur with Archit Konde
github.com/Amey-Thakur/NotebookLab

Your thinking partner, on your machine.

What “nothing leaves your machine” costs

One sentence on a landing page. Six decisions somebody has to keep making.

LOCAL API
Binds 127.0.0.1 and nothing else

AUTH TOKEN
Regenerated every launch, never written to disk

INFERENCE SIDECAR
Pinned to one release, SHA256 verified before it runs

MODEL DOWNLOADS
Restricted to HTTPS huggingface.co, in code not in docs

WEBVIEW CSP
Allows the app’s own origin, so a document cannot call home

SQL
Parameterized everywhere, because the files came from elsewhere

The claim is one line in a README. The property is six files, and it has to be chosen again every time.

Amey Thakur with Archit Konde
github.com/Amey-Thakur/NotebookLab

Your thinking partner, on your machine.

GETTING STARTED

From download to first answer

1. INSTALL

Download the build for your platform from the releases page. Windows gets an installer and an msi, macOS a dmg for both architectures, Linux a deb and an rpm. Every artifact is signed, and the app verifies an update before it installs it.

macOS is not yet notarised, so Gatekeeper will ask. Right-click, then Open, and it will let you through.

2. PICK A MODEL

On first run it suggests a local model sized to your machine. On a laptop with 16 GB and no GPU, something in the 3 to 4B instruct range answers while you are still watching. Reasoning models are worth the wait only when you can give them a GPU.

3. IMPORT

Use the picker, or open a file with NotebookLab from your file manager. Scanned pages go through OCR on the way in, so a photographed handout is searchable alongside everything typed.

4. ASK

Ask a question whose answer is a specific claim rather than a summary. Specific questions produce specific citations, which is where the tool earns its keep.

IF SOMETHING IS SLOW

A local model on a CPU is thinking, not hung. Pick a smaller one in Models, or add a cloud key and let the two share the work.

WHICH MODEL FITS YOUR MACHINE

8 GB, no GPU	16 GB, no GPU	Any GPU	Not sure
A 3B instruct model. Answers while you watch.	3 to 4B instruct. The default suggestion.	7B and up, including reasoning models.	Take the first-run suggestion. It reads your hardware.

UNDER THE HOOD

How it is put together

LAYER	WHAT IT IS
Shell	Tauri v2, so the window is the platform's own webview
Backend	Rust: parsing, OCR, retrieval, storage
Frontend	React 19 and TypeScript
Store	SQLite with WAL and FTS5
Local model	A bundled llama.cpp sidecar, any GGUF
OCR	ocrs on the pure-Rust rten runtime, no system dependency

Commands are async Tauri IPC into services, then repositories for data, providers for models, and parsers for formats. Each module is self-contained enough to remove.

NO VECTOR DATABASE

Embeddings live in SQLite and similarity is brute-force cosine in Rust. Over a few thousand chunks that is fast enough that a separate database would be an operational liability rather than an upgrade. Add one when the numbers say so, not before.

EVERYTHING IS CHECKABLE

The repository is MIT licensed and every claim in this handbook is implemented in code you can read. The security posture in particular is written down file by file, so it can be audited rather than believed.

WHAT HAPPENS TO A FILE YOU OPEN

Import
PDF, Word, Markdown, text, images

Parse
OCR runs on device for scans

Chunk
Carries file, heading and page

Embed
Vectors stored in SQLite

Retrieve
Full text and semantic

Cite
Back to the exact passage

THE END

Take it, fork it, tell us what breaks.

It is open in full, including the parts we got wrong first. Nothing here is held back for a paid tier, because there is not going to be one.

WHERE TO GO

> github.com/Amey-Thakur/NotebookLab

> amey-thakur.github.io/NotebookLab

WHAT WOULD HELP MOST

> A file it failed on, and what you expected instead.

> Which model you ran it with, and on what hardware.

> The question it answered badly. Those are the useful ones.

DOWNLOAD · FREE, NO ACCOUNT

Windows
.msi / .exe installer

macOS
.dmg, Apple Silicon & Intel

Linux
.deb / .rpm